Frequently Asked Questions

Table of contents

Questions

1. Basic terminology

1.1. What is the batch?

Batch is a group of clients with the same characteristics and loading behavior. Configuration file has one or more batches.

1.2. What means UAS?

UAS - user activity simulation - fetching url, sleeping, another url -sleeping cycles, simulating actual user activity.

2. Building curl-loader

2.1. Which operating systems are supported by curl-loader?

curl-loader supposed to work on linux with 2.4 and 2.6 kernels.

2.2. What are the building pre-requirements?

General C development environment with bash, gcc, make, etc on a linux machine is required.
Building pre-requirements are including:
1. openssl binaries;
2. openssl development package with include files (on debian package libssl-dev);
3. ncurses development package to run 'make menuconfig' with dialog GUI.
Adjust Makefile variables to point to the openssl headers and libraries. If you want to specify an openssl development directory with include files (e.g. crypto.h), export environment variable OPENSSLDIR with the value of that directory.
For example: $export OPENSSLDIR=the-full-path-to-the-directory
Another known issue is linidn.so, which means, that some linux distributions do have some libidn.so.11, but not libidn.so. Resolve it by creating a softlink.
Tarball of curl is included to the current release. When libcurl or curl-loader have building issues, correct them in the Makefile.

2.3. How to make (build) curl-loader?

Run the following commands from your (hopefully) bash linux shell:

%tar zxfv curl-loader-<version>.tar.gz
%cd curl-loader-<version>
%make
By default, we are building both libcurl and curl-loader without optimization and with debugging -g option. To build with optimization and without debugging, please, run:
%make cleanall
%make optimize=1 debug=0
If still any building issues, please, fill you free to contact us for assistance.

3. Creating Loading Configuration

3.1. How can I create loading configuration file?

To run your load, you need to create your configuration file to be passed to curl-loader by the -f commmand line option, e.g.
#curl-loader -f ./conf-user/user_batch.conf
For more examples, please, look at the files in "conf-examples" directory. You may copy an example file and edit it by using the next FAQ guidelines.

You may start with running "%make menuconfig" configuration GUI, which requires ncurses development package on your system. The dialog window will guide you and create your configuration file in conf-user directory. The current limitation of the menu-guided configuration, is that enables to create a single UAS URL.
If you need more that a single UAS url, copy your batch file from the conf-user directory to some other place and edit it by adding as much url tag series (below) as you need, but don't forget to update UAS_URLS_NUM accordingly.
The UAS url tag series:
UAS_URL=
UAS_URL_USERNAME=
UAS_URL_PASSWORD=
UAS_URL_MAX_TIME =
UAS_URL_INTERLEAVE_TIME =
For more details you may look at the next FAQ guidelines.

3.2. What are the loading configuration file tags and semantics?

Configuration file or "batch configuration file" consists from tag=value strings, grouped into 4 sections:
- general;
-login;
-UAS (user activity simulation by fetching urls and waiting for timeouts);

-logoff;
Configuration file should possess at least one client batch defined with the following params in each batch:

```
----------------------------------------------------------------------------
########### GENERAL SECTION ###################################
BATCH_NAME= bulk_batch # The name of the batch. Logfile - bulk_batch.log
CLIENTS_NUM=300 # Number of clients in the batch
CLIENTS_INITIAL_INC=30 # Clients to be added each second till CLIENTS_NUM
INTERFACE = eth0 # Name of the network interface from which to load
NETMASK=255.255.240.0 # Netmask either as an IPv4 dotted string or as a CIDR number
IP_ADDR_MIN= 192.168.1.1 # The client addresses starting address
IP_ADDR_MAX= 192.168.5.255 # Redundant, for self-control
CYCLES_NUM= 100 # Number of loading cycles to run, -1 means forever
USER_AGENT="" # User-Agent HTTP header. When empty or missed - MSIE-6 default
########### LOGIN SECTION ###################################
LOGIN=n # If 'y' or 'Y', login enabled, all other tags of the
# section to be filled. If 'n' or 'N' - comment out others
#LOGIN_USERNAME= # Redundant tag, use LOGIN_URL_USERNAME instead
#LOGIN_PASSWORD= # Redundant tag, use LOGIN_URL_PASSWORD instead
#LOGIN_REQ_TYPE= # Use string either GET+POST, POST or GET. Another option
# is to use LOGIN_REQ_TYPE_GET_POST, LOGIN_REQ_TYPE_POST or
# LOGIN_REQ_TYPE_GET tags by using tag=y to enable
#LOGIN_POST_STR= # POST string matrix. See below:
#
# To generate multiple unique users with unique passwords, use the string like
# "user=%s%d&password=%s%d". First '%s' will be substituted by the
# value of LOGIN_USERNAME tag and '%d' by the client number. Second '%s' will
# be substituted by LOGIN_PASSWORD tag value and second '%d' by the same client
# number. For example, if LOGIN_USERNAME=robert, LOGIN_PASSWORD=stam
# and LOGIN_POST_STR "user=%s%d&password=%s%d", the final POST string,
# used for the client number 1, will be "user=robert1 password;=stam1".
# In this case LOGIN_USERNAME and LOGIN_PASSWORD strings are used just
# as base-words for generating unique user credentials by appending an number.
# # To use the username and password 'as as', just provide LOGIN_POST_STR without
# %d symbols, e.g. "user=%s&secret=%s". Thus, all clients will have the same
# POST credentials with the string looking like "user=robert& secret=stam".
#
# Note, that the words like 'username', 'user', 'password', 'secret', etc are
# those fields, that login users are required to fill in their POST page.
#LOGIN_URL= # A valid http or https url to be used for login
```

```
#LOGIN_URL_USERNAME="" # Username for login url
#LOGIN_URL_PASSWORD="" # Password for login url
#LOGIN_URL_MAX_TIME= # Maximum batch time in seconds to login
#LOGIN_URL_INTERLEAVE_TIME= # Time in msec to sleep after login
#LOGIN_CYCLING= # If 'y' login should be run in cycles, and not
# just done only once
########### UAS SECTION ######################################
UAS=y # If 'y' or 'Y', login enabled, and other lines of the section to be filled
UAS_URLS_NUM = 2 # Number of urls
UAS_URL=ftp://anonymous:stam@localhost/curl-7.16.0.tar.gz # Rather large file
#UAS_URL_USERNAME="" # Username for this particular UAS url
#UAS_URL_PASSWORD="" # Password for this particular UAS url
UAS_URL_MAX_TIME = 20 # Maximum batch time in seconds to fetch the url
UAS_URL_INTERLEAVE_TIME = 0 # Time in msec to sleep after fetching the url
UAS_URL= http://localhost/apache2-default/index.html
UAS_URL_MAX_TIME = 4 # Maximum batch time in seconds to fetch the url
UAS_URL_INTERLEAVE_TIME = 0 # Time in msec to sleep after fetching the url
# You may add any number of urls providing 5-tags for each url as above,
# but do not forget to update the UAS_URLS_NUM.
########### LOGOFF SECTION ####################################
LOGOFF=n # If 'y' or 'Y', login enabled, and other tags of the
# section to be filled. If 'n' or 'N' - comment out others
#LOGOFF_REQ_TYPE= # Use strings to be GET, GET+POST, or POST
#Another option is LOGOFF_REQ_TYPE_GET_POST, LOGOFF_REQ_TYPE_POST
#or LOGOFF_REQ_TYPE_GET tags by using tag=y to enable
#LOGOFF_POST_STR= # String to be used for logoff, like "op=logoff"
#LOGOFF_URL= # A valid http or https url to be used for logoff
#LOGOFF_URL_USERNAME="" # Username for logoff url
#LOGOFF_URL_PASSWORD="" # Password for logoff url
#LOGOFF_URL_MAX_TIME= # Maximum batch time in seconds to logoff
#LOGOFF_URL_INTERLEAVE_TIME= # Time in msec to sleep after logoff
#LOGOFF_CYCLING= # If 'y' login should be run in cycles, and not just done only once
```

Worth to mention, that each batch configuration should contain all tags from the section GENERAL as well as LOGIN, UAS and LOGOFF section tags. When LOGIN, UAS or LOGOFF is set as 'y' (yes), all tags for that section should appear (uncommented) and to be set to some valid values (empty string "" to be used not to define LOGIN_POST_STR and LOGOFF_POST_STR tags, when POST is not applicable). In the case, that LOGIN, UAS or LOGOFF section tag is disabled by setting 'n' (no) value, thus, all the tags of the disabled section may appear with with empty strings "", without values or just may be commented out by '#'.

TAG CLIENTS_INITIAL_INC:
It serves for gradual increase of clients number at the loading initial phase. Use the tag in GENERAL section to specify the number of loading clients to be added each second till the total clients number reaches the number specified by CLIENTS_NUM tag.

Note, that both quotted and non-quotted string are supported as the tags values.
For more examples, please, look at the files in "conf-examples" directory.

3.3. How does the configuration support login, logoff and authentication flavors?

curl-loader preforms Login and logoff operations using the following HTTP methods:
- GET+POST (server response to GET provides a post-form to be filled and posted by POST);
- POST only;
- GET only.
UAS urls similar to Login urls are coming with an option to configure username and password. The difference is that only GET method is currently allowed in UAS mode and username+&password are solely to support Web or Proxy Authentication (see below). Login url uniquely allows POSTing user credentials via configurable POST forms as well as configuring a unique username and password for each virtual client by adding a sequence number to the user and password basewords.
Does somebody really needs PUT or POST methods as an option for UAS urls?
If yes, please, signal us.
The loader supports HTTP Web Authentication and Proxy Authentication. Authentication methods supported are Basic, Digest (RFC2617) and NTLM. When responded 401 or 407, libcurl will choose the most safe method from those, supported by the server.
To support GSS Web-Authentication, add in Makefile building of libcurl against appropriate GSS library, see libcurl pages for detailed instructions.

3.4. What about FTP load?

To generate FTP/FTPS load, please, use UAS section and pass user credentials via ftp-url according to the RFC 1738 like:
ftp://username:password@hostname:port/etc-str
Please, look at conf-examples/ftp.conf

4. Running load

4.1. What are the running environment requirements?

Running hundreds and thousands of clients, please, do not forget:

- to increase limit of descriptors (sockets) by running e.g.
#ulimit -n 10000;
- optionally, to set reuse of sockets in time-wait state: by setting
#echo 1 > /proc/sys/net/ipv4/tcp_tw_recycle and/or
#echo 1 > /proc/sys/net/ipv4/tcp_tw_reuse;
In some cases you may need to increase the system limits for open descriptors (sockets).

4.2. How I can run the load?

Usage: run as a root user:
#./curl-loader -f <configuration filename> [other options]
Other possible options are:
-c[onnection establishment timeout, seconds]
-e[rror drop client (smooth mode). Client on error doesn't attempt loading any mode]
-h[elp]
-i[ntermediate (snapshot) statistics time interval (default 3 sec)]
-f[ilename of configuration to run (batches of clients)]
-l[ogfile max size in MB (default 1024). On the size reached, file pointer rewinded]
-m[ode of loading, 1 - storming, 2 - smooth (default)]
-o[utput to stdout bodies of downloaded files - attn!- bulky]
-r[euse connections disabled. Close connections and re-open them. Try with and without]
-s[tderr printout of client messages instead of to logfile - attn!- bulky]
-t[hreads enable - enables threads, each runs a batch of clients]
-v[erbose output to the logfiles; includes info about headers sent/received]
-u[rl logging - logs url names to logfile, when -v verbose option is used]
For the rare cases, when several batches are specified in the same config file, note, please, that curl-loader without -t runs only the first batch specified. Thus, running several client batches without threads requires some script with several curl-loader processes.
Connection Reuse Disable Option (-r):
The default behavior of curl-loader after HTTP response is to re-use the tcp-connection for the next request. If you are specifying -r command-line option, the TCP connection will be closed and re-opened for the next request. Whether it is appropriate for you to work with -r option or without, it depends on your server support of Keep-Alive and the purpose of your testing. Try with and without -r and see, what you get.
Connection reuse (just the default, without -r option) has advantages due to the decreased consumption of opened descriptors (sockets) and ports.

4.3. Which loading modes are supported?

One of the modes is called "storming" (-m1 command line), where all clients of a batch are starting together. They are expected to finish their job within a certain timeout. After the

timeout they are either completed fetching url, or are cutted and disconnected and a new cycle begins. Storming mode is better to use, when login operation is a single one and not in cycles, or when the tool is used just as a traffic generator.

Another mode is a so-called "smooth" (the default mode or -m2 command line), where each client spends as much time as wishes and starts another url or another cycle only after completing the previous url. When a client loading is terminated due to some reason, e.g. connection timeout, the default behavior is to schedule loading for this client at the next cycle. If -e option is passed to command line, the client will not be re-scheduled any more, which is usefull to get indication of errors by monitoring drop in number of active clients. Look in the logfile for errors, and when connection timeout error appears, adjust the connection timeout using -c command line option. Note, that the smooth mode is suitable for login-logoff cycles of load.

A combination of the both modes (two separate instances of the program) can provide somehow realistic behavior with some "smooth" load plus "storming" bursts of requests. Hyper-mode (option -m3) is currently "under construction".

4.4. How I can monitor loading progress status?

curl-loader outputs to the console loading status and statistics as the "standing" ouput, where the upper part is for the latest interval and below is the average numbers since load start. The example is below:

```
=========================================================
Last interval stats (interval:3 sec, clients:10, CAPS:3):
Operations: Success Failed
LOGIN: 4 1
UAS-0: 3 0
LOGOFF: 2 0
HTTP/FTP-Req:10,2xx:5,3xx:4,4xx:0,5xx:1,Err:0,Delay:20,Delay-2xx:12,Thr-in:12015(b/s),Thr-out:35511(
HTTPS/FTPS-Req:0,2xx:0,3xx:0,4xx:0,5xx:0,Err:0,Delay:0,Delay-2xx:0,Thr-in:0(b/s),Thr-out:0(b/s)
-------------------------------------------------------
Summary stats since load start (load runs:24 secs, CAPS-average:4):
Operations: Success Failed
LOGIN: 0 32
UAS-0: 0 0
LOGOFF: 0 0
HTTP/FTP-Req:95,2xx:47,3xx:48,4xx:0,5xx:3,Err:1,Delay:31,Delay-2xx:15,Thr-in:113921(b/s),Thr-out:33(
HTTPS/FTPS-Req:0,2xx:0,3xx:0,4xx:0,5xx:0,Err:0,Delay:0,Delay-2xx:0,Thr-in:0(b/s),Thr-out:0(b/s)
=========================================================
```

A copy of the output is also saved in the file <batch-name>.txt

4.5. Where is the detaled log of all virtual clients activities and how to read it?

DETAILED LOGFILE is written to the file named:
<batch-name>.log:
The semantics of logfile output, using command line options -v (verbous) and -u (url print):
"Cycle number", "Client number (ip-address)" - some information string, e.g.:
4 39 (192.168.0.39) :== Info: Trying 10.30.6.42... : eff-url:
http://10.30.6.42:8888/server/Admin/ServiceList.do, url:
Which meas: cycle: 4, client number 39 with ipv4 address (192.168.0.39), status of the
message is Info, eff-url - is the url, used right now, "url:" is empty, which means, that it is the
same as effective.
Effective url may be a result of redirection and, thus, "url:"
(target url, specified in batch configuration file) will be printed as well.
Please, note, that when the logfile reaches 1024 MB size, curl-loader rewinds it and starts to
overwrite it from the beginning. Y may tune the rewinding file size by using command line
option:
-l <log-filesize-in-MB>

4.6. Which statistics is collected and how to get to it?

Currently HTTP/HTTPS statistics includes the following counters:
- requests num;
- 2xx success num;
- 3xx redirects num;
- client 4xx errors num;
- server 5xx errors num;
- other errors num, like resolving, tcp-connect, server closing or empty responses number;
- average application server delay (msec), estimated as the time between HTTP request and
HTTP response without taking into the account network latency (RTT);
- average application server delay for 2xx (success) HTTP-responses, as above, but only for
2xx responses. The motivation for that is that 3xx redirections and 5xx server errors/rejects
may not necessarily provide a true indication of a testing server working functionality.
- throughput out (batch average);
- throughput in (batch average);
The statistics goes to the screen (both the interval and the current summary statistics for the
load) as well as to the file with name <batch_name>.txt When the load completes or when
the user presses CTRL-C (sometimes some clients may stall), the final load report is printed
at the console as well as to the statistics file.
Some strings from the file:

Run-Time,Appl,Clients,Req,2xx,3xx,4xx,5xx,Err,Delay,Delay-2xx,Thr-In,Thr-Out
2, HTTP/FTP, 100, 155, 0, 96, 0, 0, 1154, 1154, 2108414, 15538
2, HTTPS/FTPS, 100, 0, 0, 0, 0, 0, 0, 0, 0, 0
4, HTTP/FTP, 100, 75, 32, 69, 0, 0, 1267, 1559, 1634656, 8181
4, HTTPS/FTPS, 100, 0, 0, 0, 0, 0, 0, 0, 0, 0
Cutted here
36, HTTP/FTP, 39, 98, 35, 58, 0, 0, 869, 851, 1339168, 11392
36, HTTPS/FTPS, 39, 0, 0, 0, 0, 0, 0, 0, 0, 0
38, HTTP/FTP, 3, 91, 44, 62, 0, 0, 530, 587, 1353899, 10136
38, HTTPS/FTPS, 3, 0, 0, 0, 0, 0, 0, 0, 0, 0
*, *, *, *, *, *, *, *, *, *, *, *, *
Run-Time,Appl,Clients,Req,2xx,3xx,4xx,5xx,Err,Delay,Delay-2xx,Thr-In,Thr-Out
38, HTTP/FTP, 0, 2050, 643, 1407, 0, 213, 725, 812, 1610688, 11706
38, HTTPS/FTPS, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0

The bottom strings after asterisks are for final averages.
At the same time a clients dump file with name <batch_name>.ctx is generated to provide detailed statistics about each client state and statistics counters.
One string from the file:
1 (192.168.0.1)
,cycles:124,cstate:1,b-in:22722029,b-out:174605,req:745,2xx:497,3xx:248,4xx:0,5xx:0,err:0
where
1 (192.168.0.1)- is the index of the client and its ip-address;
cycles- number of loading cycles done;
cstate - is the number of the client state (-1 - error, 0 - init, 1- login, 2- uas, 3-logoff, 4-final-ok);
b-in - bytes passed in;
b-out - bytes passed out;
req- number of requests done;
2xx, 3xx, 4xx, 5xx - number of responses Nxx received;
err - number of libcurl errors at the resolving, TCP/IP and TLS/SSL levels;
The following conditions are considered as errors:
- error at the level of libcurl, which includes resolving, TCP/IP and, when applicable, TLS/SSL errors;
- all HTTP 5xx server errors;
- most of HTTP 4xx client errors, excluding are 401 and 407 authentication responses not considered as errors;
When the above error conditions occur, a virtual client is marked as being in the error state. By default we "recover" such client by scheduling it to the next loading cycle, starting from the first operation of the cycle. You may use command line option -e to change the default

behavior to another, so that clients once arriving at error state will not be scheduled for load any more.

5. Advanced Issues

5.1. What about performance?

The limit of no more than 1000 virtual clients per batch (due to the select () FD_SETSIZE) has been lifted, so you can try now 2000 and may be more number of clients from the single thread (single select ()). Please, care however about available sockets and limits. Note, that the system call select () does not scale well with a big number of sockets, therefore, you can try a script, which runs several instances of curl_loader, each with 1000-2000 clients. Running several batches with threads (option -t), each with 1000-2000 clients is another valid option, however, may be less stable - try it.
A future release will incorporate libcurl HYPER API to support tens of thousand of virtual clients from a single curl-loader process, using curl_multi_socket() approach and epoll API support by libevent as in http://curl.haxx.se/lxr/source/hiper/hipev.c

5.2. How to run a really big load?

These are the temporary instructions till we hopefully migrate to epoll-based I/O demultiplexing (or even kevent based?), supporting tens of thousand of clients in one thread without select/poll limitations.
1. Compile with optimization;
Since you need performance compile with optimization and without debugging.
%make cleanall
%make optimize=1 debug=0
Y may add to Makefile optimization for your particular processor by -match /-mcpu gcc option directives to OPT_FLAGS.
2. Login as a su;
3. Increase the default number of allowed open descriptors (sockets);
E.g. running #ulimit -n 20000
If running several instances of curl-loader, consider increase of system limits for open descriptors, if necessary. Take your own account of the socket usage in the system, considering sockets faster recycling (less time in the time-wait state), by setting, optionally, something like this:
#echo 1 > /proc/sys/net/ipv4/tcp_tw_recycle and/or
#echo 1 > /proc/sys/net/ipv4/tcp_tw_reuse;
Correct, if required, the value of CURL_LOADER_FD_SETSIZE (set to 20000 in Makefile) to control the maximum fd, that may be used for select.

Increase the maximum number of open descriptors in your linux system, if required, using linux HOWTOS.

4. Create configuration files for each instance of curl-loader to run.

What is important is to give a unique BATCH_NAME for each batch-file, which is in use by a separate instance of curl-loader. Logfile, report file, etc have name, which are derivatives of the BATCH_NAME, and when several instances of curl-loader are writing to the same file, this is not helpful and crashful. Please, use in the configuration batch files non-overlapping ranges of IP-addresses, else libcurl virtual clients will compete for the IP-addresses to bind to.

Use CLIENTS_INITIAL_INC tag in smooth mode to increase number of your clients gradually at start-up in order not boom the server.

5. Connections re-use.

The default behavior of curl-loader now is after HTTP response to re-use the tcp-connection for the next request-response. If you are specifying -r command-line option, the connection will be closed and re-opened for the next request. Whether it is appropriate for you to work with -r or without depends on your server support of Keep-Alive and the purpose of your testing.

Try with and without -r and see, what you get.

6. Try sometimes the non-default storming mode (commandline option -m1), which may better suit your needs. See, README for explanations. A combination of -m1 and -r is a valid combination, still, if the server does not make problems with connections reuse.

7. Troubleshooting.

Running first runs one may try command-line options -v (verbose) and -u (url in logs). Grep to look for the errors and their reasons. If an error is "Connection timeout", you may try to increase the connection establishment timeout (the default is 5 seconds - huge, but "ih veis"), using -c command-line option.

If any assistance required, please, don't hesitate to contact us.

8. Logs and statistics.

After end of a run, or after SIGINT (Cntl-C), the final results are calculated and printed to the console as well as to the file <batch-name>.txt. Current results are presented in each row, and average summary as the last raws, separated from the rest by asterisks.

Pay attention, that <batch-name>.log log file may become huge, particularly, when using verbose output (-v -u). Commandline option -l <maxsize in MB> may be useful, whereas the default policy is to rewind the logfile (writing from the file start), when it reaches 1 GB. Otherwise, do not use -v and -u options for rather long loads.

5.3. How to calculate CAPS numbers for a load?

When number of clients is defined by CLIENTS_NUM tag, the number of CAPS (call attempts per seconds) is resulting from the clients number and load duration for each cycle, comprising from:

- login time with possible redirections and sleeping after login interval;
- uas time for each url with possible redirections, intervals betwee urls and after uas interval;
- logoff time with possible redirections and sleeping after logoff interval;

The actions and time intervals are configurable in batch file, whereas url retrival time is server and network dependent and not always easy to predict. The result is that number of clients/requests is a known parameter, and number of CAPS is something to be estimated from the time of test and number of requests.

Smooth mode presents at the LOAD STATUS GUI output calculated current and average CAPS numbers.